

Images, Video, and Media

Media is where content quality and Core Web Vitals collide, and where most sites quietly fail both.

Bottom line

Images are usually the largest thing on a page and the most common cause of a failing Largest Contentful Paint. Getting them right serves accessibility, search visibility, and performance at the same time. There are about six rules and they are all cheap.

The weight problem

Your hero image is very often your LCP element, which means it directly determines whether the page passes the threshold you will meet properly in Section 27.

Targets:

Guideline	Value	Tier
Hero image weight	under 150KB	[practitioner]
Format	WebP or AVIF	[practitioner]
LCP threshold the image must not break	2.5 seconds at p75	[confirmed]

WebP is universally supported now and typically 25 to 35% smaller than an equivalent JPEG. **AVIF** is smaller still with slightly less support. Serving WebP with a JPEG fallback is the safe default.

Size the image to its display size. A 4000px image displayed at 800px wastes most of the bytes. Responsive srcset lets the browser pick.

Never lazy-load the hero image. Lazy loading below-the-fold images is good. Lazy loading the LCP element delays the exact thing being measured. Mark it `fetchpriority="high"` and load it eagerly.

Always set width and height attributes. Without them the browser does not know how much space to reserve, the layout jumps when the image arrives, and that is Cumulative Layout Shift. This one attribute pair prevents an entire category of CLS failure.

Alt text

Every non-decorative image needs descriptive alt text. It is an accessibility requirement first, a search signal second, and the two want the same thing.

Describe the image, in context. What would you say to someone who cannot see it and needs to understand why it is here?

- Bad: `alt="image"`, `alt="seo audit seo checklist seo tools"`, or empty on a meaningful image
- Good: `alt="Search Console coverage report showing 340 pages indexed and 82 excluded"`

Purely decorative images take an empty alt (`alt=""`), which tells screen readers to skip them. Omitting the attribute entirely is not the same thing.

Do not keyword-stuff. Alt text stuffed with terms is a recognizable spam pattern and it actively degrades the accessibility function it exists for.

File names

`search-console-coverage-report.webp` beats `IMG_4471.webp`. It costs nothing at upload time and is annoying to fix later.

Captions

Captions are read at roughly the same rate as headlines, far more than body text. If an image is carrying real information, caption it. Captions are also genuine content that search engines index, unlike alt text which is a fallback.

Video

Host on YouTube unless you have a reason not to. YouTube is the most-cited domain in Google AI Overviews and correlates strongly with ChatGPT visibility, because models are trained on transcripts. Self-hosting gets you control and costs you that entire discovery surface.
[practitioner]

Publish the transcript on the page. This is the highest-value and most-skipped video SEO action. The transcript is indexable text, it is what AI systems can actually retrieve, and it makes the content accessible. A video with no transcript is invisible to text retrieval.

Do not autoplay with sound. It is a page experience problem and people leave.

Embed lazily. A YouTube embed pulls a significant amount of JavaScript. Use a facade, meaning a thumbnail that loads the real player on click, unless the video is the point of the page.

Media that earns citations

From Section 5, the content types that still get clicked and cited overlap heavily with original media:

- **Original diagrams and charts** built from your own data
- **Screenshots of real interfaces** with real numbers in them
- **Comparison tables**, which extract cleanly into AI answers and featured snippets

Stock photography does none of this. A generic photo of people around a laptop adds page weight and nothing else. If an image is not carrying information, consider whether it should exist.

Why this matters

media is the most common single cause of a failing Core Web Vitals assessment, and images are usually the easiest large win available. Meanwhile alt text and transcripts are the cheapest accessibility work you can do, and they happen to be exactly what makes media legible to retrieval systems.

Do this now

- 1 **Run PageSpeed Insights on your Section 9 page.** Note the LCP element. It is probably an image.
- 2 **Check your hero image weight.** Over 150KB, compress it and convert to WebP. Squoosh is free and browser-based.
- 3 **Confirm the hero is not lazy-loaded** and has `fetchpriority="high"`.
- 4 **Add width and height attributes** to every image on the page.
- 5 **Audit alt text** on that page. Every meaningful image described, decorative ones empty, nothing stuffed.
- 6 **Rename any IMG_1234-style files** going forward. Do not bulk-rename existing ones, that is a URL change with the costs from Section 13.
- 7 **If the page has a video, publish the transcript** underneath it.
- 8 **Delete one stock image** that carries no information. Note the weight saved.
- 9 **Re-run PageSpeed Insights** and compare LCP.

Capstone step

Your priority page now has compressed modern-format images sized to display, a hero that loads eagerly with high priority, dimension attributes preventing layout shift, real alt text, and a transcript if there is video. You have a before-and-after LCP number.

Key takeaways

- The hero image is usually the LCP element. Under 150KB, WebP or AVIF, sized to display, never lazy-loaded, `fetchpriority="high"`.

- Width and height attributes on every image prevent an entire category of Cumulative Layout Shift failure.
- Alt text is accessibility first and search second, and both want the same thing: describe the image in context. Empty alt for decorative images, never stuffed.
- Publish video transcripts. Untranscribed video is invisible to text retrieval, and YouTube is the most-cited domain in AI Overviews.