

Sitemaps and Indexation

A sitemap is a claim about what deserves indexing. Fill it with junk and Google stops believing you.

Bottom line

XML sitemaps are a discovery aid, not a ranking factor. Their real value is diagnostic: a clean sitemap plus the Search Console indexing report tells you exactly which pages Google refuses to store and why. Healthy sites hold an indexed-to-submitted ratio above 85%.

What a sitemap is for

It helps discovery, especially for large sites, new sites with few backlinks, and pages that are poorly internally linked.

It is not a ranking factor. Google has confirmed this. Submitting a sitemap does not improve rankings and never has. [confirmed]

It is a statement of intent. You are telling Google "these are my canonical, indexable, valuable URLs". That statement is only useful if it is true.

Its best use is diagnostic. Submit a sitemap of exactly the pages you believe should be indexed, then compare against what actually is. The gap is your work list.

The hard limits

Limit	Value
URLs per sitemap file	50,000
Uncompressed file size	50MB
Solution beyond that	sitemap index file pointing at multiple sitemaps

[confirmed]

What belongs in it

Only URLs that are all of the following:

- Return **200 OK**
- Are **canonical**, not pointing elsewhere via a canonical tag

- Are **indexable**, no noindex
- Are **not disallowed** in robots.txt
- You genuinely **want indexed**

What must not be in it: 404s, redirects, noindexed pages, canonicalized duplicates, parameter variants, paginated pages beyond the first, tag and author archives you do not care about, thank-you and confirmation pages.

The reason this matters more than it sounds: **including junk trains Google to distrust the sitemap**. If a meaningful proportion of submitted URLs are dead or noindexed, the whole file becomes a weaker signal. [practitioner]

Most CMS-generated sitemaps include everything by default. Check yours rather than assuming.

Structure for diagnosis

Split sitemaps by content type rather than shipping one giant file:

```
`` /sitemap.xml (index) /sitemap-pages.xml /sitemap-blog.xml
/sitemap-services.xml /sitemap-images.xml ``
```

This costs nothing and turns Search Console into a much better report: you can see that 98% of service pages are indexed but only 60% of blog posts, which localizes the problem immediately.

lastmod is worth setting accurately. Google uses it as a hint about what to recrawl. Setting it to today's date on every page every day destroys its value and is a recognized bad pattern. [practitioner]

priority and changefreq are ignored. Do not spend time on them.

Reading the indexation report

Search Console, Pages. This is the report that matters.

Healthy: above 85% of submitted URLs indexed. [practitioner]

The exclusion reasons worth knowing:

Reason	What it means	Fix
Crawled - currently not indexed	Google fetched it and declined. Quality signal	Improve or remove the page
Discovered - currently not indexed	Known but not fetched. Crawl signal	Internal links, flatter architecture, crawl budget
Duplicate, Google chose different canonical	Google overruled your canonical	Section 23

Reason	What it means	Fix
Alternate page with proper canonical tag	Working as intended	Nothing, but should not be in the sitemap
Excluded by noindex	Working as intended	Nothing, but remove from sitemap
Soft 404	Returns 200 but looks empty or error-like	Return a real 404, or add content
Blocked by robots.txt	Cannot be crawled	Section 21
Page with redirect	Not a final URL	Remove from sitemap

You met the first two in Section 2. They look adjacent and share no fixes: one is a content problem, the other is an architecture problem.

Getting pages indexed faster

Request indexing in URL Inspection. It queues the URL. It does not guarantee speed and the daily quota is small, so use it for genuinely important pages, not in bulk.

Internal links are the real mechanism. A new page linked from a page Google crawls often gets found quickly. An orphan may never be found at all.

IndexNow is a push protocol: you notify search engines that a URL changed rather than waiting. Supported by **Bing, Yandex and Naver**. Google does **not** support it. Worth doing anyway, because Bing powers ChatGPT search, which makes Bing indexation strategically more important than its market share implies. [practitioner]

Most major CMSs have an IndexNow plugin. It takes ten minutes.

When pages should not be indexed

Deliberate exclusion is a legitimate tool. `noindex` these:

- Thank-you and confirmation pages
- Internal search results
- Thin tag and category archives nobody searches for
- Staging and development environments, and check this, since leaked staging sites are common
- Paginated pages beyond the first, in most cases
- Filtered and faceted URL variants

Remember from Section 21: use `noindex`, not `Disallow`, or Google cannot read the instruction.

Why this matters

indexation is binary. A page that is not indexed cannot rank, and no amount of content or link work changes that. It is also the fastest technical audit you can run, because Search Console tells you the answer directly and most people never open the report.

Do this now

- 1 Find your sitemap.** Usually `/sitemap.xml` or `/sitemap_index.xml`. Confirm it loads.
- 2 Spot-check 10 URLs from it.** Do they all return 200, are they canonical, are any noindexed?
- 3 Remove anything that should not be there.** Most CMS sitemaps include junk by default.
- 4 Split by content type** if you have more than a few hundred URLs.
- 5 Submit it in Search Console** and confirm it is read without errors.
- 6 Open the Pages report** and calculate your indexed-to-submitted ratio.
- 7 Record the counts for "Crawled - not indexed" and "Discovered - not indexed".** Compare against the numbers you wrote down in Section 2.
- 8 Pick three "Crawled - not indexed" pages and read them honestly.** These are quality rejections.
- 9 Pick three "Discovered - not indexed" pages and check their internal links.** These are usually orphans or buried too deep.
- 10 Install IndexNow** if your CMS supports it.

Capstone step

Your sitemap contains only canonical, indexable, 200-status URLs, split by content type. You know your indexed-to-submitted ratio and have separated the quality rejections from the crawl-reach problems, with three examples of each identified for fixing.

Key takeaways

- Sitemaps aid discovery and are not a ranking factor. Their real value is diagnostic.
- Include only 200-status, canonical, indexable URLs you actually want indexed. Junk in the sitemap trains Google to distrust it.
- Above 85% indexed-to-submitted is healthy. Split sitemaps by content type so the report localizes the problem for you.
- IndexNow covers Bing, Yandex and Naver but not Google. Do it anyway, because Bing powers ChatGPT search.