

Canonicals and Duplicate Content

There is no duplicate content penalty. There is signal dilution, which is worse because nobody tells you it is happening.

Bottom line

Duplicate content does not get you penalized. It splits ranking signals across several URLs so that none of them is as strong as the one page you should have had. The canonical tag is a hint, not a directive, and Search Console tells you when Google overrules you.

The penalty myth

There is no duplicate content penalty for ordinary duplication. Google's position is consistent: duplication is normal on the web, and it deduplicates rather than punishes. [confirmed]

What actually happens is worse in practice because it is silent:

- **Signals split.** Links, engagement and relevance spread across several URLs.
- **Google picks one** to represent the group, and it may not be the one you wanted.
- **Crawl budget is wasted** fetching the same content repeatedly.

Penalties apply to *scraped* or deliberately manipulative duplication, which is a different thing.

Where duplication comes from

Most duplication is accidental and technical.

Protocol and host variants. `http://` and `https://`, `www` and `non-www`. Four versions of your homepage if unresolved.

Trailing slashes. `/services` and `/services/`.

Case sensitivity. `/Services` and `/services` on servers that treat them as distinct.

URL parameters. Tracking (`?utm_source=`), sorting (`?sort=price`), filtering, and session IDs all produce unique URLs for identical content.

Pagination. Page 2 of a category often shares most of its template and description with page 1.

Faceted navigation. The largest generator of duplicates on ecommerce sites. Colour x size x brand x price produces thousands of near-identical URLs.

Printer-friendly and AMP versions.

Syndicated content. Your article republished elsewhere, or someone else's on your site.

Boilerplate-heavy pages. Twenty location pages differing only in the city name are functionally duplicates, and this one is a *content* problem wearing a technical costume.

The canonical tag

```
`html <link rel="canonical" href="https://example.com/preferred-url" /> `
```

It is a hint, not a directive. Google considers it alongside other signals and can pick a different canonical. This surprises people who expect it to be binding. [confirmed]

Use a self-referencing canonical on every indexable page. The page points at itself. This is the default and it protects against parameter variants appearing.

Rules that make canonicals work:

- Use **absolute URLs**, not relative
- Point at a URL that returns **200**, not a redirect and not a 404
- **One canonical per page.** Multiple tags cause Google to ignore all of them
- **Do not canonicalize everything to the homepage.** A common and destructive error that tells Google the rest of your site does not exist
- **Canonical and noindex together is contradictory.** One says "index this other page instead", the other says "index nothing". Pick one
- **Internal links should point at canonical URLs directly**, not at variants that canonicalize elsewhere. Contradicting your own canonical with your internal linking weakens it [practitioner]

When Google overrules you

Search Console, URL Inspection, shows "**Google-selected canonical**" next to "**User-declared canonical**". When they differ, Google disagreed.

Common reasons:

- The canonical target is substantially different content, so the pages are not actually duplicates
- Internal links overwhelmingly point at the other URL
- The other URL is in the sitemap and this one is not
- The other version has the backlinks
- Redirect and canonical signals contradict each other

The fix is to make every signal agree: canonical tag, internal links, sitemap inclusion, and redirects all pointing the same way.

The right tool for each case

Situation	Tool
Same content, one permanent preferred URL	301 redirect
Same content, both URLs must remain accessible	Canonical tag
Near-duplicate you never want indexed	<code>noindex</code>
Infinite parameter combinations	Canonical plus robots.txt on the crawl trap
Content genuinely serving different intents	Rewrite so they are different
Same content, different languages or regions	hreflang , Section 31

Redirect when you can, canonicalize when you must. A 301 is a stronger, unambiguous signal. Canonicals are for when both URLs need to stay reachable.

Parameters and faceted navigation

The biggest source of duplication at scale.

- **Self-referencing canonical on the clean URL**, and canonical from parameter variants back to it.
- **`noindex` on filter combinations** nobody searches for.
- **Robots-disallow genuine crawl traps**, meaning infinite combinations, accepting from Section 21 that disallowed URLs cannot be deindexed by a `noindex` they can never read. Clean them up first, then block.
- **Keep internal links pointing at clean URLs.** If your own faceted nav links to parameterized URLs, you are generating the problem yourself.

The location-page trap

Worth calling out because agencies do it constantly.

Twenty pages that are identical except for a city name are not twenty pages. Google will index one or two and ignore the rest, or index none. This is not solved with canonicals, because they are not duplicates you want consolidated, they are pages you want to be genuinely distinct.

The fix is content: local specifics, real projects in that area, genuine differences. Section 35 covers this properly. If you cannot write a genuinely different page for a location, you probably should not have one.

Why this matters

duplication is silent. Nothing in Search Console shouts about it, no penalty notice arrives, and the pages still exist. They just quietly underperform, and the cause is invisible unless you go looking at which canonical Google actually selected.

Do this now

- 1 Test your host and protocol variants.** Load `http://`, `https://`, `www` and `non-www`. All should redirect to one canonical version.
- 2 Test trailing slash behaviour** on a few URLs.
- 3 Run URL Inspection on 5 important pages.** Compare user-declared canonical to Google-selected canonical. Note any disagreement.
- 4 Check for self-referencing canonicals** on your key templates.
- 5 Search for parameter URLs in the index:** `site:yourdomain.com inurl:?`. Note what is indexed that should not be.
- 6 Verify no page has multiple canonical tags**, which is a common plugin conflict.
- 7 Confirm nothing is canonicalized to the homepage** except the homepage.
- 8 Check internal links point at canonical URLs**, not at variants.
- 9 If you have location or service-area pages, read two side by side.** If they differ only by a place name, mark them for the Section 35 rewrite.

Capstone step

Host, protocol and trailing-slash variants all resolve to one version. Self-referencing canonicals are in place, no page carries conflicting signals, and you know of any case where Google has overruled your canonical choice.

Key takeaways

- There is no duplicate content penalty. There is signal dilution, which is worse because it is silent and nothing reports it.
- The canonical tag is a hint. Search Console's "Google-selected canonical" tells you when you were overruled, and the fix is making every signal agree.
- Redirect when you can, canonicalize when both URLs must stay reachable, `noindex` when you simply never want it indexed.
- Location pages differing only by a city name are a content problem, not a canonical problem.