

How AI Search Works

One prompt becomes ten to twenty searches, and the winner is decided before your page is ever read.

Bottom line

AI search does not rank pages, it retrieves passages and synthesizes an answer. A single prompt fans out into **10 to 20 sub-queries**, Google narrows **200 to 500 candidate documents down to 5 to 15 citations**, and the overlap between ranking in the organic top 10 and being cited has collapsed from **92% in mid-2025 to roughly 38% in early 2026**. Ranking first no longer implies being quoted.

The mechanism, in one pass

Traditional search: query goes in, an ordered list of ten pages comes out, you click one.

AI search runs a **multi-stage pipeline**:

- 1 Interpretation.** The system works out what you actually mean, including implicit sub-questions you did not ask.
- 2 Query fan-out.** Your prompt is expanded into multiple synthetic sub-queries run in parallel.
- 3 Retrieval.** Each sub-query pulls candidate documents from an index.
- 4 Selection.** Candidates are narrowed hard. **200 to 500 documents down to 5 to 15 cited sources.** [practitioner]
- 5 Synthesis.** A model writes one answer from the selected passages and attributes some of them.

The important structural point: you are not competing for a position, you are competing to be selected as source material. Everything in Sections 37 through 40 follows from that.

Query fan-out

The single most important concept in this tier.

A prompt expands into **5 to 11 sub-queries** typically, **10 to 20** on deeper analysis, and **hundreds of parallel searches** for deep-research modes. [practitioner]

Ask an engine "what's the best CRM for a small agency" and behind the scenes it may run: what is a CRM, best CRM for small business, CRM pricing comparison, CRM for agencies specifically, HubSpot vs Pipedrive, CRM features small teams need, CRM reviews 2026, and a dozen more.

You do not need to win the original query. You need to be the best answer to several of the sub-queries. That is a completely different optimization target than a keyword ranking, and it is why keyword-level thinking under-performs here.

AI queries are also far longer: 70 to 80 words on average versus 3 to 4 for a traditional search, a **17 to 26x complexity increase**. People type full situations into AI, not keywords. [practitioner]

The coverage finding

This is the most actionable number in the entire AI-search block.

Sites with 80%+ topical coverage of their domain retain 85.4% of AI visibility despite fan-out instability. [practitioner]

The reasoning is direct. Fan-out is unpredictable: the same prompt generates different sub-queries on different runs. A site with one strong page on a topic wins only when a sub-query happens to match that page. A site that covers the whole topic area comprehensively gets caught by whichever sub-queries fire.

Depth of coverage is the durable strategy. Individual page optimization is not. This is the Section 15 topical cluster argument, restated with an AI-specific number attached, and it is the reason clusters were worth building before AI search existed and are worth more now.

The overlap collapse

Period	Organic top-10 to AI citation overlap
Mid-2025	92%
Early 2026	~38%
AI Mode specifically	14 to 17%

[practitioner]

This is the number that justifies AI search being its own discipline rather than a footnote to SEO.

In 2025, optimizing for Google essentially optimized for AI citation as a side effect. That is no longer true. Roughly six in ten cited sources are now pages that do not rank in the organic top ten for the query being answered.

The practical read: your ranking report and your citation report are now two different reports measuring two different things, and a client can be winning one while losing the other. Section 41 handles measuring both.

Do not overcorrect. Ranking still matters, both for direct traffic and for entering the candidate pool. The collapse means AI visibility needs its own work, not that classic SEO stopped mattering.

Retrieval favors passages, not pages

Because synthesis works from retrieved passages, the unit of competition is smaller than a page.

The corpus is specific: content should be structured as **self-contained answer units of roughly 134 to 167 words**. [practitioner]

An answer unit is a chunk that makes sense in isolation. It states what it is about, answers one thing, and does not depend on the paragraph above it for context. A section that opens "As we discussed earlier, this approach..." is unusable as a retrieved passage because extracted alone it means nothing.

This is why Sections 12 and 14 emphasized standalone section structure. Section 37 turns it into a concrete rewriting method.

What this changes about strategy

Traditional SEO	AI search
Compete for a position	Compete to be selected as source material
Optimize a page for a query	Cover a topic well enough to catch many sub-queries
Ranking is the outcome	Citation is the outcome, and ranking does not predict it
The page is the unit	The passage is the unit
Links are the authority signal	Mentions correlate ~3x more strongly (Section 34)

Both columns are live. You are running two overlapping games, not replacing one with the other.

Why this matters

most people encountering AI search reach for a tactic, usually schema or llms.txt, because a tactic is easy to sell and easy to implement. The mechanism tells you why those are minor. Selection happens across many fan-out sub-queries from a pool of hundreds of candidates, which rewards comprehensive topical coverage and clear extractable passages far more than any markup.

Do this now

- 1 **Pick your most important commercial query.**
- 2 **Ask ChatGPT, Perplexity and Google AI Mode the same question**, phrased as a real person would in 70 words, not as a keyword.
- 3 **Write down every source each one cited.** Note how much they overlap. They usually do not.

- 4 Check whether the cited pages rank in Google's top 10** for your query. Most will not. That is the 38% collapse in front of you.
- 5 List the sub-questions each answer covered.** That is observable fan-out, and it is your real content target list.
- 6 Score your own topical coverage.** Of the sub-questions you just listed, how many do you have a genuinely good page for? That fraction is your coverage estimate against the 80% benchmark.
- 7 Open your best page and check its passages.** Does any 150-word chunk stand alone and answer one thing? If every section depends on the one above it, note that for Section 37.
- 8 Record which platform cited you, if any.** This is your citation baseline.

Capstone step

You have traced one core query's fan-out across three engines, recorded the cited sources and whether they rank organically, listed the observable sub-questions as a content target list, and scored your capstone site's topical coverage against the 80% benchmark.

Key takeaways

- AI search retrieves passages and synthesizes. You compete to be selected as source material, not to hold a position.
- Query fan-out expands one prompt into **10 to 20 sub-queries**. Win several sub-queries, not the original one.
- **80%+ topical coverage retains 85.4% of AI visibility.** Comprehensive clusters beat optimized individual pages because fan-out is unpredictable.
- Organic top-10 to citation overlap fell from **92% to ~38%**. Ranking and citation are now two different reports, and you need both.